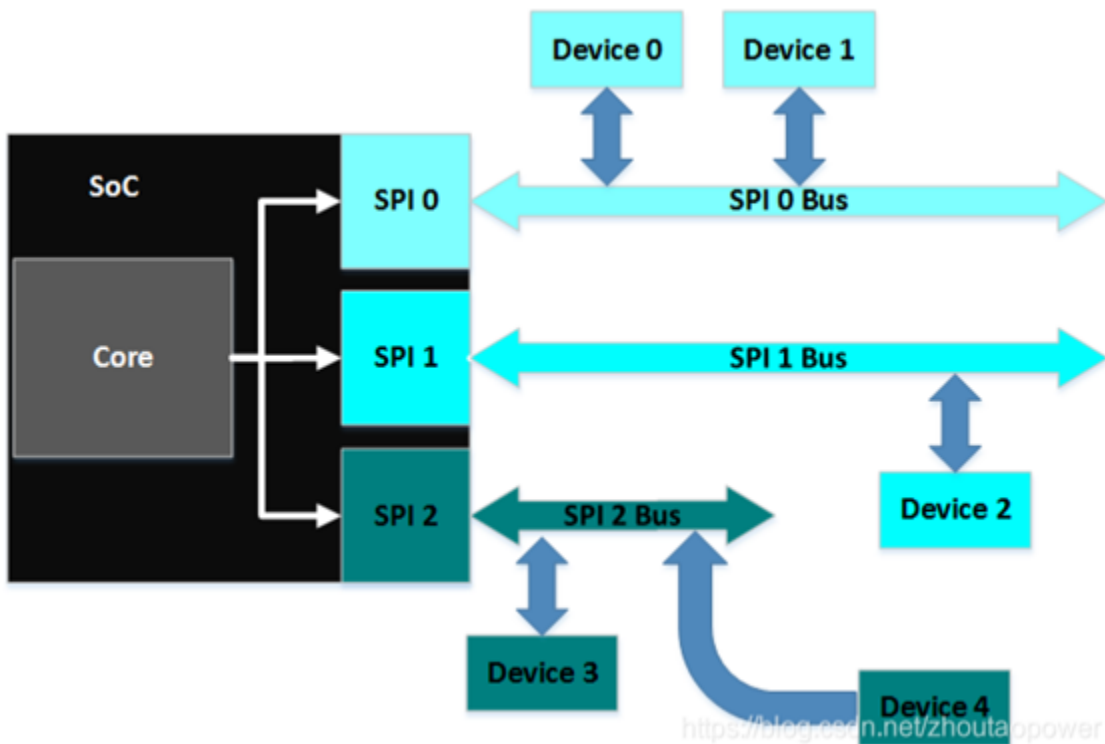


SPI(serial peripheral interface)

SPI是一种高速的串行全双工的接口，在SoC中被广泛的使用。在linux内核中，SPI和我们较为熟悉的platform一样，属于总线的一类，并且拥有自己的一套框架，该框架大致将SPI的驱动分为了以下几个部分：

1. SPI主机：即SPI controller，老版本中叫做SPI master
2. SPI外设驱动：外设指的是SPI controller上的从设备
3. SPI 从设备描述
4. SPI传输层的动作描述

SPI硬件拓扑



一般来说SoC上都会提供多个SPI控制器，每个SPI控制器上又可以挂载多个从设备。

内核SPI硬件描述数据结构

根据SPI的硬件拓扑结构，linux内核中的SPI框架将其抽象为主设备（spi_controller），从设备（spi_device）和对应的驱动（spi_driver）。

spi_controller

```
struct spi_controller {
    struct device      dev;

    struct list_head list;

    /* other than negative (== assign one dynamically), bus_num is fully
     * board-specific. usually that simplifies to being SOC-specific.
     * example: one SOC has three SPI controllers, numbered 0..2,
     * and one board's schematics might show it using SPI-2. software
     * would normally use bus_num=2 for that controller.
     */
    s16      bus_num;

    /* chipselects will be integral to many controllers; some others
     * might use board-specific GPIOs.
     */
    u16      num_chipselect;

    ...

    /* flag indicating this is an SPI slave controller */
    bool      slave;

    ...

    int      (*setup)(struct spi_device *spi);

    /*
     * set_cs_timing() method is for SPI controllers that supports
     * configuring CS timing.
     *
     * This hook allows SPI client drivers to request SPI controllers
     * to configure specific CS timing through spi_set_cs_timing() after
     * spi_setup().
     */
    int (*set_cs_timing)(struct spi_device *spi, struct spi_delay *setup,
                        struct spi_delay *hold, struct spi_delay *inactive);

    /* bidirectional bulk transfers
     *
     * + The transfer() method may not sleep; its main role is
     *   just to add the message to the queue.
     * + For now there's no remove-from-queue operation, or
     *   any other request management
     * + To a given spi_device, message queueing is pure fifo
     *
     * + The controller's main job is to process its message queue,
     *   selecting a chip (for masters), then transferring data
     * + If there are multiple spi_device children, the i/o queue
     *   arbitration algorithm is unspecified (round robin, fifo,
     *   priority, reservations, preemption, etc)
     *
     * + Chipselect stays active during the entire message
     *   (unless modified by spi_transfer.cs_change != 0).
     * + The message transfers use clock and SPI mode parameters
     *   previously established by setup() for this device
     */
    int      (*transfer)(struct spi_device *spi,
```

```

        struct spi_message *mesg);

/* called on release() to free memory provided by spi_controller */
void                                (*cleanup)(struct spi_device *spi);

/*
 * Used to enable core support for DMA handling, if can_dma()
 * exists and returns true then the transfer will be mapped
 * prior to transfer_one() being called. The driver should
 * not modify or store xfer and dma_tx and dma_rx must be set
 * while the device is prepared.
 */
...
bool                                queued;
struct kthread_worker                *kworker;
struct kthread_work                  pump_messages;
spinlock_t                           queue_lock;
struct list_head                     queue;
...
int (*prepare_transfer_hardware)(struct spi_controller *ctlr);
int (*transfer_one_message)(struct spi_controller *ctlr,
                            struct spi_message *mesg);
int (*unprepare_transfer_hardware)(struct spi_controller *ctlr);
int (*prepare_message)(struct spi_controller *ctlr,
                       struct spi_message *message);
int (*unprepare_message)(struct spi_controller *ctlr,
                         struct spi_message *message);
int (*slave_abort)(struct spi_controller *ctlr);

/*
 * These hooks are for drivers that use a generic implementation
 * of transfer_one_message() provided by the core.
 */
void (*set_cs)(struct spi_device *spi, bool enable);
int (*transfer_one)(struct spi_controller *ctlr, struct spi_device *spi,
                   struct spi_transfer *transfer);
void (*handle_err)(struct spi_controller *ctlr,
                   struct spi_message *message);

...
};

```

该结构体用来表示一个spi控制器，完整的数据结构可以在/include/linux/spi/spi.h中查看，这里保留了比较重要的成员，包括：

```

bus_numspi spi bus
num_chipselect: spi
mode_bitsSPI slave mode
slaveslavespi
probesetuptransferset_cs
kworkerSPI

```

spi_device

```

struct spi_device {
    struct device          dev;
    struct spi_controller *controller;
    struct spi_controller *master;      /* compatibility layer */
    u32                    max_speed_hz;
    u8                      chip_select;
    u8                      bits_per_word;
    bool                    rt;
    u32                     mode;

#define SPI_CPHA            0x01          /* clock phase */
#define SPI_CPOL            0x02          /* clock polarity */
#define SPI_MODE_0          (0|0)        /* (original MicroWire) */
#define SPI_MODE_1          (0|SPI_CPHA)
#define SPI_MODE_2          (SPI_CPOL|0)
#define SPI_MODE_3          (SPI_CPOL|SPI_CPHA)
#define SPI_CS_HIGH         0x04          /* chipselect active high? */
#define SPI_LSB_FIRST       0x08          /* per-word bits-on-wire */
#define SPI_3WIRE           0x10          /* SI/SO signals shared */
#define SPI_LOOP            0x20          /* loopback mode */
#define SPI_NO_CS           0x40          /* 1 dev/bus, no chipselect */
#define SPI_READY           0x80          /* slave pulls low to pause */
#define SPI_TX_DUAL         0x100        /* transmit with 2 wires */
#define SPI_TX_QUAD         0x200        /* transmit with 4 wires */
#define SPI_RX_DUAL         0x400        /* receive with 2 wires */
#define SPI_RX_QUAD         0x800        /* receive with 4 wires */
#define SPI_CS_WORD         0x1000       /* toggle cs after each word */
#define SPI_TX_OCTAL        0x2000       /* transmit with 8 wires */
#define SPI_RX_OCTAL        0x4000       /* receive with 8 wires */
#define SPI_3WIRE_HIZ       0x8000       /* high impedance turnaround */

    int                     irq;
    void                    *controller_state;
    void                    *controller_data;
    char                    modalias[SPI_NAME_SIZE];
    const char              *driver_override;
    int                     cs_gpio;      /* LEGACY: chip select gpio */
    struct gpio_desc         *cs_gpiod;   /* chip select gpio desc */
    struct spi_delay         word_delay; /* inter-word delay */

    /* the statistics */
    struct spi_statistics    statistics;

    /*
     * likely need more hooks for more protocol options affecting how
     * the controller talks to each chip, like:
     * - memory packing (12 bit samples into low bits, others zeroed)
     * - priority
     * - chipselect delays
     * - ...
     */
};

```

从设备的信息较少，主要包括了：

```

controller:
chip_selectnum_chipselect
modeSPI
modalias

```

spi_driver

```

/**
 * struct spi_driver - Host side "protocol" driver
 * @id_table: List of SPI devices supported by this driver
 * @probe: Binds this driver to the spi device. Drivers can verify
 *         that the device is actually present, and may need to configure
 *         characteristics (such as bits_per_word) which weren't needed for
 *         the initial configuration done during system setup.
 * @remove: Unbinds this driver from the spi device
 * @shutdown: Standard shutdown callback used during system state
 *            transitions such as powerdown/halt and kexec
 * @driver: SPI device drivers should initialize the name and owner
 *          field of this structure.
 *
 * This represents the kind of device driver that uses SPI messages to
 * interact with the hardware at the other end of a SPI link. It's called
 * a "protocol" driver because it works through messages rather than talking
 * directly to SPI hardware (which is what the underlying SPI controller
 * driver does to pass those messages). These protocols are defined in the
 * specification for the device(s) supported by the driver.
 *
 * As a rule, those device protocols represent the lowest level interface
 * supported by a driver, and it will support upper level interfaces too.
 * Examples of such upper levels include frameworks like MTD, networking,
 * MMC, RTC, filesystem character device nodes, and hardware monitoring.
 */
struct spi_driver {
    const struct spi_device_id *id_table;
    int (*probe)(struct spi_device *spi);
    int (*remove)(struct spi_device *spi);
    void (*shutdown)(struct spi_device *spi);
    struct device_driver driver;
};

```

spi_driver中只有三个函数指针，其注释已经说明，该驱动作用于从设备。

spi_board_info

在内核的spi框架中，同样提供了板级启动时使用的从设备结构体，既然是描述从设备的，那么该结构体中的内容和spi_device几乎没有差别。

```

struct spi_board_info {
    /* the device name and module name are coupled, like platform_bus;
     * "modalias" is normally the driver name.
     */
    /* platform_data goes to spi_device.dev.platform_data,
     * controller_data goes to spi_device.controller_data,
     * device properties are copied and attached to spi_device,
     * irq is copied too
     */
    char                modalias[SPI_NAME_SIZE];
    const void          *platform_data;
    const struct property_entry *properties;
    void                *controller_data;
    int                 irq;

    /* slower signaling on noisy or low voltage boards */
    u32                 max_speed_hz;

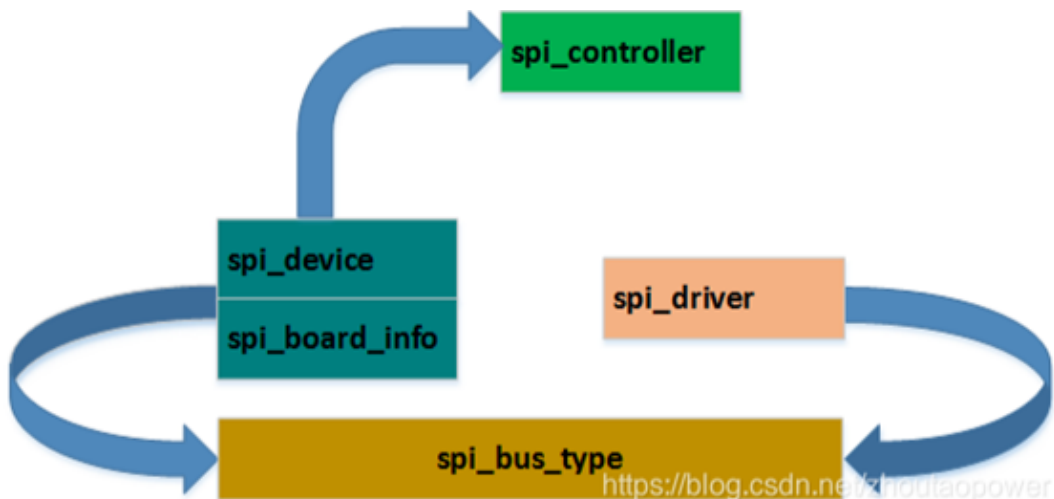
    /* bus_num is board specific and matches the bus_num of some
     * spi_controller that will probably be registered later.
     */
    /* chip_select reflects how this chip is wired to that master;
     * it's less than num_chipselect.
     */
    u16                 bus_num;
    u16                 chip_select;

    /* mode becomes spi_device.mode, and is essential for chips
     * where the default of SPI_CS_HIGH = 0 is wrong.
     */
    u32                 mode;

    /* ... may need additional spi_device chip config data here.
     * avoid stuff protocol drivers can set; but include stuff
     * needed to behave without being bound to a driver:
     * - quirks like clock rate mattering when not selected
     */
};

```

内核中硬件描述关系总结：



SPI信息交换机制分析

在介绍完linux内核中关于SPI硬件部分的抽象后，接下来就到了描述SPI信息交换的环节。这部分主要涉及了两个结构体，`spi_transfer`和`spi_message`。

`spi_transfer`

```
struct spi_transfer {
    /* it's ok if tx_buf == rx_buf (right?)
     * for MicroWire, one buffer must be null
     * buffers must work with dma_*map_single() calls, unless
     * spi_message.is_dma_mapped reports a pre-existing mapping
     */
    const void      *tx_buf;
    void            *rx_buf;
    unsigned        len;

    dma_addr_t      tx_dma;
    dma_addr_t      rx_dma;
    struct sg_table tx_sg;
    struct sg_table rx_sg;

    unsigned        cs_change:1;
    unsigned        tx_nbits:3;
    unsigned        rx_nbits:3;
#define SPI_NBITS_SINGLE      0x01 /* 1bit transfer */
#define SPI_NBITS_DUAL       0x02 /* 2bits transfer */
#define SPI_NBITS_QUAD       0x04 /* 4bits transfer */
    u8              bits_per_word;
    u16             delay_usecs;
    struct spi_delay delay;
    struct spi_delay cs_change_delay;
    struct spi_delay word_delay;
    u32             speed_hz;

    u32             effective_speed_hz;

    unsigned int    ptp_sts_word_pre;
    unsigned int    ptp_sts_word_post;

    struct ptp_system_timestamp *ptp_sts;

    bool            timestamped;

    struct list_head transfer_list;

#define SPI_TRANS_FAIL_NO_START BIT(0)
    u16             error;
};
```

`spi_transfer`是spi消息传递的最小单位，在`spi_transfer`中，定义了这个传输中所需的发送与接收的内存buffer（`rx/tx_buf`），速率，传输长度，和所需的`delay`等。其他的则是在启用`dma`功能时所需要的信息。

多个`spi_transfer`构成了`spi_message`，在注释中，`spi_message`被描述为一个完整的多片的`spi`传输。

spi_message

```
struct spi_message {
    struct list_head    transfers;

    struct spi_device    *spi;

    unsigned            is_dma_mapped:1;

    /* REVISIT: we might want a flag affecting the behavior of the
     * last transfer ... allowing things like "read 16 bit length L"
     * immediately followed by "read L bytes". Basically imposing
     * a specific message scheduling algorithm.
     *
     * Some controller drivers (message-at-a-time queue processing)
     * could provide that as their default scheduling algorithm. But
     * others (with multi-message pipelines) could need a flag to
     * tell them about such special cases.
     */

    /* completion is reported through a callback */
    void                (*complete)(void *context);
    void                *context;
    unsigned            frame_length;
    unsigned            actual_length;
    int                status;

    /* for optional use by whatever driver currently owns the
     * spi_message ... between calls to spi_async and then later
     * complete(), that's the spi_controller controller driver.
     */
    struct list_head    queue;
    void                *state;

    /* list of spi_res resources when the spi message is processed */
    struct list_head    resources;
};
```

其成员变量主要就由一个`spi_transfer`构成的list，和`spi`异步消息交换时需要的回调。

接下来我们关注一下`spi`消息交换的大致流程。

spi消息交换流程分析

我们只关注spi内核框架中关于消息交换相关的函数，首先最重要的就是spi控制器初始化时需要适配的底层最终的transfer函数，该函数共有三个变种，在最初的spi框架中只有transfer，而如今的spi控制器则较多的去实现了后面两个变种，transfer_one和transfer_one_message。

transfer_one函数指挥硬件完成一个spi_transfer消息的交换，而transfer_one_message可以指挥硬件完成一个spi_message的交换。

观察spi_register_controller函数中关于transfer钩子相关的部分：

spi_register_controller

```
/*
 * If we're using a queued driver, start the queue. Note that we don't
 * need the queueing logic if the driver is only supporting high-level
 * memory operations.
 */
if (ctlr->transfer) {
    dev_info(dev, "controller is unqueued, this is deprecated\n");
} else if (ctlr->transfer_one || ctlr->transfer_one_message) {
    status = spi_controller_initialize_queue(ctlr);
    if (status) {
        device_del(&ctlr->dev);
        goto free_bus_id;
    }
}
/* add statistics */
```

可以看到，transfer作为老式的spi控制器如今只为了兼容老机器所使用，新的spi控制器都会选择实现transfer_one或者transfer_one_message（c2000中的spi控制器实现了transfer_one）。之后内核会初始化spi控制器中的队列，跟进该函数：

spi_controller_initialize_queue

```
static int spi_controller_initialize_queue(struct spi_controller *ctlr)
{
    int ret;

    ctlr->transfer = spi_queued_transfer; //
    if (!ctlr->transfer_one_message)
        ctlr->transfer_one_message = spi_transfer_one_message; //

    /* Initialize and start queue */
    ret = spi_init_queue(ctlr); //
    if (ret) {
        dev_err(&ctlr->dev, "problem initializing queue\n");
        goto err_init_queue;
    }
    ctlr->queued = true;
    ret = spi_start_queue(ctlr); //
    if (ret) {
        dev_err(&ctlr->dev, "problem starting queue\n");
        goto err_start_queue;
    }

    return 0;

err_start_queue:
    spi_destroy_queue(ctlr);
err_init_queue:
    return ret;
}
```

可以看到，该函数将transfer和transfer_one_message替换为了内核的函数，然后初始化并启动了该spi控制器的消息队列。

init&start queue

```
static int spi_init_queue(struct spi_controller *ctlr)
{
    ctlr->running = false;
    ctlr->busy = false;

    ctlr->kworker = kthread_create_worker(0, dev_name(&ctlr->dev)); // worker
    if (IS_ERR(ctlr->kworker)) {
        dev_err(&ctlr->dev, "failed to create message pump kworker\n");
        return PTR_ERR(ctlr->kworker);
    }

    kthread_init_work(&ctlr->pump_messages, spi_pump_messages); // workspi_pump_messages

    /*
     * Controller config will indicate if this controller should run the
     * message pump with high (realtime) priority to reduce the transfer
     * latency on the bus by minimising the delay between a transfer
     * request and the scheduling of the message pump thread. Without this
     * setting the message pump thread will remain at default priority.
     */
    if (ctlr->rt)
        spi_set_thread_rt(ctlr);

    return 0;
}

static int spi_start_queue(struct spi_controller *ctlr)
{
    unsigned long flags;

    spin_lock_irqsave(&ctlr->queue_lock, flags);

    if (ctlr->running || ctlr->busy) {
        spin_unlock_irqrestore(&ctlr->queue_lock, flags);
        return -EBUSY;
    }

    ctlr->running = true;
    ctlr->cur_msg = NULL;
    spin_unlock_irqrestore(&ctlr->queue_lock, flags);

    kthread_queue_work(ctlr->kworker, &ctlr->pump_messages); //workerworkspi_pump_messages

    return 0;
}
```

记住spi_pump_messages这个函数。我们介绍过spi在linux内核中支持消息的同步（spi_sync）和异步（spi_async）传输，其差异在于同步传输会等到spi_message传输完成才返回，我们接下来看一看spi异步传输的机制，看下spi_pump_messages这个函数是如何在其中发挥作用的。

__spi_async

```
static int __spi_async(struct spi_device *spi, struct spi_message *message)
{
    struct spi_controller *ctlr = spi->controller;
    struct spi_transfer *xfer;

    /*
     * Some controllers do not support doing regular SPI transfers. Return
     * ENOTSUPP when this is the case.
     */
    if (!ctlr->transfer)
        return -ENOTSUPP;

    message->spi = spi; // messagespi

    SPI_STATISTICS_INCREMENT_FIELD(&ctlr->statistics, spi_async);
    SPI_STATISTICS_INCREMENT_FIELD(&spi->statistics, spi_async);

    trace_spi_message_submit(message);

    if (!ctlr->ptp_sts_supported) {
        list_for_each_entry(xfer, &message->transfers, transfer_list) {
            xfer->ptp_sts_word_pre = 0;
            ptp_read_system_prets(xfer->ptp_sts);
        }
    }

    return ctlr->transfer(spi, message); // controllertransferspi
}
```

__spi_queued_transfer

```
static int __spi_queued_transfer(struct spi_device *spi,
                                struct spi_message *msg,
                                bool need_pump)
{
    struct spi_controller *ctlr = spi->controller;
    unsigned long flags;

    spin_lock_irqsave(&ctlr->queue_lock, flags);

    if (!ctlr->running) {
        spin_unlock_irqrestore(&ctlr->queue_lock, flags);
        return -ESHUTDOWN;
    }
    msg->actual_length = 0;
    msg->status = -EINPROGRESS;

    list_add_tail(&msg->queue, &ctlr->queue); // controller
    if (!ctlr->busy && need_pump)
        kthread_queue_work(ctlr->kworker, &ctlr->pump_messages); // workerspi_pump_messages

    spin_unlock_irqrestore(&ctlr->queue_lock, flags);
    return 0;
}
```

可以看到，`spi_message`的传输最终都会在`spi_pump_message`中进行。`spi_pump_message`的函数整体比较长，大致的流程包括：

1. 检查当前是否有msg在传输

2. 检查当前controller是否为繁忙状态

3. 调用底层的prepare_message函数，进行一些传输前的设置

4. 如果开启了dma功能则进行一些相应的设置

4. 最后调用controller上的函数transfer_one_message，如果spi控制器没有实现该函数，则该函数在之前的流程中被设置为了

spi_transfer_one_message

spi_transfer_one_message

```
static int spi_transfer_one_message(struct spi_controller *ctrlr,
                                   struct spi_message *msg)
{
    struct spi_transfer *xfer;
    bool keep_cs = false;
    int ret = 0;
    struct spi_statistics *statm = &ctrlr->statistics;
    struct spi_statistics *stats = &msg->spi->statistics;

    spi_set_cs(msg->spi, true, false);

    SPI_STATISTICS_INCREMENT_FIELD(statm, messages);
    SPI_STATISTICS_INCREMENT_FIELD(stats, messages);

    list_for_each_entry(xfer, &msg->transfers, transfer_list) { // msgtransferstransfermsg
        trace_spi_transfer_start(msg, xfer);

        spi_statistics_add_transfer_stats(statm, xfer, ctrlr);
        spi_statistics_add_transfer_stats(stats, xfer, ctrlr);

        if (!ctrlr->ptp_sts_supported) {
            xfer->ptp_sts_word_pre = 0;
            ptp_read_system_prets(xfer->ptp_sts);
        }

        if ((xfer->tx_buf || xfer->rx_buf) && xfer->len) {
            reinit_completion(&ctrlr->xfer_completion);
        }
    }

    fallback_pio:

    ret = ctrlr->transfer_one(ctrlr, msg->spi, xfer); // transfer_one
    if (ret < 0) {
        if (ctrlr->cur_msg_mapped &&
            (xfer->error & SPI_TRANS_FAIL_NO_START)) {
            __spi_unmap_msg(ctrlr, msg);
            ctrlr->fallback = true;
            xfer->error &= ~SPI_TRANS_FAIL_NO_START;
            goto fallback_pio;
        }

        SPI_STATISTICS_INCREMENT_FIELD(statm,
                                       errors);
        SPI_STATISTICS_INCREMENT_FIELD(stats,
                                       errors);

        dev_err(&msg->spi->dev,
                "SPI transfer failed: %d\n", ret);
        goto out;
    }

    if (ret > 0) {
        ret = spi_transfer_wait(ctrlr, msg, xfer); // spi
        if (ret < 0)
            msg->status = ret;
    }
}
```

```

        } else {
            if (xfer->len)
                dev_err(&msg->spi->dev,
                        "Bufferless transfer has length %u\n",
                        xfer->len);
        }

        if (!ctlr->ptp_sts_supported) {
            ptp_read_system_postts(xfer->ptp_sts);
            xfer->ptp_sts_word_post = xfer->len;
        }

        trace_spi_transfer_stop(msg, xfer);

        if (msg->status != -EINPROGRESS)
            goto out;

        spi_transfer_delay_exec(xfer);

        if (xfer->cs_change) {
            if (list_is_last(&xfer->transfer_list,
                            &msg->transfers)) {
                keep_cs = true;
            } else {
                spi_set_cs(msg->spi, false, false);
                _spi_transfer_cs_change_delay(msg, xfer);
                spi_set_cs(msg->spi, true, false);
            }
        }

        msg->actual_length += xfer->len;
    }

out:
    if (ret != 0 || !keep_cs)
        spi_set_cs(msg->spi, false, false);

    if (msg->status == -EINPROGRESS)
        msg->status = ret;

    if (msg->status && ctlr->handle_err)
        ctlr->handle_err(ctlr, msg);

    spi_finalize_current_message(ctlr); // controller

    return ret;
}

```

至此，linux内核驱动控制spi控制器成功的完成了一次消息的交换。

linux内核中SPI驱动框架的大体结构如下：

